

Chapter 7

October 30, 2019

```
In [1]: %reload_ext sql
```

```
In [2]: %sql postgresql://postgres:postgres@localhost:5432/analysis
```

```
Out[2]: 'Connected: postgres@analysis'
```

1 Table Design that Works for You

1.1 Primary Key Syntax

```
In [3]: %%sql
```

```
CREATE TABLE natural_key_example (  
    licence_id varchar(10) CONSTRAINT license_key PRIMARY KEY,  
    first_name varchar(10),  
    last_name varchar(50)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

```
Out[3]: []
```

- we first use the column constraint syntax to declare **licence_id** as the primary key
- followed by a name for the primary key and then the keywords **PRIMARY KEY**
- an advantage of this is it is easy to understand at a glance which column is the primary key

```
In [4]: %%sql
```

```
DROP TABLE natural_key_example;
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

```
Out[4]: []
```

- we simply delete the table

In [5]: %%sql

```
CREATE TABLE natural_key_example (  
    licence_id varchar(10),  
    first_name varchar(10),  
    last_name varchar(50),  
    CONSTRAINT license_key PRIMARY KEY (licence_id)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

Out[5]: []

- here we use the CONSTRAINT after listing the final column
- we also give here the column which we want to declare as the key after giving the name of the key and the keyword PRIMARY KEY
- if you want to use more than one primary key you must use this syntax
- you then declare each primary key with commas (**composite primary key**)

In [6]: %%sql

```
INSERT INTO natural_key_example (licence_id,  
                                first_name, last_name)  
VALUES ('T229901', 'Lynn', 'Malero');
```

```
* postgresql://postgres:***@localhost:5432/analysis  
1 rows affected.
```

Out[6]: []

- we insert some data into our created table

In [7]: %%sql

```
SELECT * FROM natural_key_example;
```

```
* postgresql://postgres:***@localhost:5432/analysis  
1 rows affected.
```

Out[7]: [('T229901', 'Lynn', 'Malero')]

- if we do the same operation (inserting some values into our table)
- since the id has to be unique, we will get an error
- a unique primary key protects the table from runining the integrity of the data

1.2 Creating a Composite Primary Key

In [10]: `%%sql`

```
CREATE TABLE natural_key_composite_example (  
    student_id varchar(10),  
    school_day date,  
    present boolean,  
    CONSTRAINT student_key PRIMARY KEY (student_id, school_day)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

Out[10]: []

- if we want to create a composition of multiple columns as one primary key we this with the *composite primary key*
- the combination of unique student id and date column the student was in school each day during a school year
- the present colum of the type boolean represents if the student was there (present) in that day

1.3 Creating an Auto-Incrementing Surrogate Key

In [11]: `%%sql`

```
CREATE TABLE surrogate_key_example (  
    order_number bigserial,  
    product_name varchar(50),  
    order_date date,  
    CONSTRAINT order_key PRIMARY KEY (order_number)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

Out[11]: []

In [12]: `%%sql`

```
INSERT INTO surrogate_key_example (product_name, order_date)  
VALUES ('Beachball Polish', '2015-03-17'),  
      ('Wrinkle De-Atomizer', '2017-05-22'),  
      ('Flux Capacitor', '1985-10-26');
```

```
* postgresql://postgres:***@localhost:5432/analysis  
3 rows affected.
```

Out[12]: []

In [13]: %%sql

```
SELECT * FROM surrogate_key_example;
```

```
* postgresql://postgres:***@localhost:5432/analysis
3 rows affected.
```

Out[13]: [(1, 'Beachball Polish', datetime.date(2015, 3, 17)),
(2, 'Wrinkle De-Atomizer', datetime.date(2017, 5, 22)),
(3, 'Flux Capacitor', datetime.date(1985, 10, 26))]

- we used bigserial for autoincrementing the primary key
- when you insert data into the table you can omit the order_number column
- it will be added automatically and incremented also automatically

1.3.1 Foreign Keys

In [34]: %%sql

```
CREATE TABLE licenses (  
    license_id varchar(10),  
    first_name varchar(50),  
    last_name varchar(50),  
    CONSTRAINT licenses_key PRIMARY KEY (license_id)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out[34]: []

In [35]: %%sql

```
CREATE TABLE registrations (  
    registration_id varchar(10),  
    registration_date date,  
    license_id varchar(10) REFERENCES licenses (license_id),  
    CONSTRAINT registration_key PRIMARY KEY  
    (registration_id, license_id)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out[35]: []

In [36]: %%sql

```
INSERT INTO licenses (license_id, first_name, last_name)
VALUES ('T229901', 'Lynn', 'Malero');
```

```
* postgresql://postgres:***@localhost:5432/analysis
1 rows affected.
```

Out[36]: []

In []: %%sql

```
INSERT INTO registrations (registration_id, registration_date, license_id)
VALUES ('A203391', '3/17/2017', 'T229901');
```

In []: %%sql

```
INSERT INTO registrations (registration_id, registration_date, license_id)
VALUES ('A75772', '3/17/2017', 'T000001');
```

- the last insert statement had an licence_id which was not in the licences table
- our foreign key was not existed so we get an error
- this is somehow good, because we can guarantee the integrity of the data

1.3.2 Automatically Deleting Related Records with CASCADE

- to delete a row in licenses and have that action automatically delete any related rows in registrations
 - we can specify this with
 - * **ON DELETE CASCADE**

1.3.3 The CHECK Constraint

In [43]: %%sql

```
CREATE TABLE check_constraint_example (
    user_id bigserial,
    user_role varchar(50),
    salary integer,
    CONSTRAINT user_id_key PRIMARY KEY (user_id),
    CONSTRAINT check_role_in_list CHECK (user_role IN('Admin', 'Staff')),
    CONSTRAINT check_salary_not_zero CHECK (salary > 0)
);
```

```
* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out [43]: []

- a **CHECK** constraint evaluates whether the data added to a column meets the expected criteria which we specify with a logical test
- if this criteria aren't met, we get an error
- it can prevent columns having nonsensical data (such as a grade 'Z')
 - **CHECK (logical expression)** after all columns are defined
- we can also combine more than one logical test with **AND**

1.3.4 The UNIQUE Constraint

In [49]: %%sql

```
CREATE TABLE unique_constraint_example (  
    contact_id bigserial CONSTRAINT contact_id_key PRIMARY KEY,  
    first_name varchar(50),  
    last_name varchar(50),  
    email varchar(200),  
    CONSTRAINT email_unique UNIQUE (email)  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

Out [49]: []

In [50]: %%sql

```
INSERT INTO unique_constraint_example (first_name, last_name, email)  
VALUES ('Samantha', 'Lee', 'slee@example.org');
```

```
* postgresql://postgres:***@localhost:5432/analysis  
1 rows affected.
```

Out [50]: []

In [51]: %%sql

```
INSERT INTO unique_constraint_example (first_name, last_name, email)  
VALUES ('Betty', 'Diaz', 'bdiaz@example.org');
```

```
* postgresql://postgres:***@localhost:5432/analysis  
1 rows affected.
```

Out [51]: []

In []: %%sql

```
INSERT INTO unique_constraint_example (first_name, last_name, email)
VALUES ('Sasha', 'Lee', 'slee@example.org');
```

- this will occur an error
- the email column we expect this to be unique
- those addresses change over time and can be null
- we use **UNIQUE** to ensure to have one email address for one contact
- the main difference between this and the primary key is, that this value can be null

1.3.5 The NOT NULL Constraint

In [55]: %%sql

```
CREATE TABLE not_null_example (
    student_id bigserial,
    first_name varchar(50) NOT NULL,
    last_name varchar(50) NOT NULL,
    CONSTRAINT student_id_key PRIMARY KEY (student_id)
);
```

```
* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out [55]: []

- Sometimes we want columns not be null
- we declare this with the **NOT NULL** keyword
- if we attend an insert for the table and don't include values for those columns, the database will notify us the violation

1.3.6 Removing Constraints or Adding Them Later

- to remove a primary key, foreign key, or a unique constraint you would use this statement:
 - ALTER TABLE table_name DROP CONSTRAINT constraint_name;
- to drop a NOT NULL constraint, we have to use the ALTER COLUMN statement also:
 - ALTER TABLE table_name ALTER COLUMN column_name DROP NOT NULL;

In [60]: %%sql

```
ALTER TABLE not_null_example DROP CONSTRAINT student_id_key;
```

```
* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out [60]: []

- we drop the primary key

In [62]: %%sql

```
ALTER TABLE not_null_example ADD CONSTRAINT student_id_key PRIMARY KEY (student_id);

* postgresql://postgres:***@localhost:5432/analysis
(psycopg2.errors.InvalidTableDefinition) multiple primary keys for table "not_null_example" are
[SQL: ALTER TABLE not_null_example ADD CONSTRAINT student_id_key PRIMARY KEY (student_id);]
(Background on this error at: http://sqlalche.me/e/f405)
```

- we add a new primary key

In [64]: %%sql

```
ALTER TABLE not_null_example ALTER COLUMN first_name DROP NOT NULL;

* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out [64]: []

- we drop a column which is NOT NULL

In [59]: %%sql

```
ALTER TABLE not_null_example ALTER COLUMN first_name SET NOT NULL;

* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out [59]: []

- we set a new column which is NOT NULL

1.4 Speeding Up Queries with Indexes

- In the same way a book indexes, SQL indexes are working the same
- the primary key can be seen as an Index too
- there are many other index types which will be handled now

1.4.1 B-Tree: PostgreSQL's Default Index

In [65]: %%sql

```
CREATE TABLE new_york_adresses (  
    longitude numeric(9,6),  
    latitude numeric(9,6),  
    street_number varchar(10),  
    street varchar(32),  
    unit varchar(7),  
    postcode varchar(5),  
    id integer CONSTRAINT new_york_key PRIMARY KEY  
);
```

```
* postgresql://postgres:***@localhost:5432/analysis  
Done.
```

Out [65]: []

In []: %%sql

```
COPY new_york_adresses  
FROM '/Users/ugurtigu/Documents/Learn/Docs/SQL/city_of_new_york.csv'  
WITH (FORMAT CSV, HEADER);
```

In [77]: %%sql

```
SELECT * FROM new_york_adresses  
LIMIT 10;
```

```
* postgresql://postgres:***@localhost:5432/analysis  
10 rows affected.
```

Out [77]: [(Decimal('-73.939546'), Decimal('40.725332'), '608', 'MORGAN AVENUE', None, '11222',
(Decimal('-73.923303'), Decimal('40.692213'), '962', 'BUSHWICK AVENUE', None, '11221',
(Decimal('-73.918562'), Decimal('40.700381'), '309', 'HARMAN STREET', None, '11237',
(Decimal('-73.923921'), Decimal('40.693764'), '1115', 'GREENE AVENUE', None, '11221',
(Decimal('-73.918152'), Decimal('40.699605'), '1341', 'GREENE AVENUE', None, '11237',
(Decimal('-73.916312'), Decimal('40.701468'), '1415', 'GREENE AVENUE', None, '11237',
(Decimal('-73.937288'), Decimal('40.704293'), '185', 'MOORE STREET', None, '11206', 3
(Decimal('-73.938597'), Decimal('40.704417'), '3', 'BUSHWICK COURT', None, '11206', 3
(Decimal('-73.941511'), Decimal('40.703775'), '100', 'MOORE STREET', None, '11206', 3
(Decimal('-73.960914'), Decimal('40.704585'), '135', 'RODNEY STREET', None, '11211',

Benchmarking Query Performance with Explain

- We will measure the performance of the query with the EXPLAIN keyword

In [82]: %%sql

```
EXPLAIN ANALYSE SELECT * FROM new_york_addresses
WHERE street = 'BROADWAY';
```

```
* postgresql://postgres:***@localhost:5432/analysis
7 rows affected.
```

```
Out[82]: [('Bitmap Heap Scan on new_york_addresses (cost=76.23..6368.64 rows=3072 width=46) (actual t
("  Recheck Cond: ((street)::text = 'BROADWAY'::text)",),
('  Heap Blocks: exact=2157',),
('   ->  Bitmap Index Scan on street_idx (cost=0.00..75.46 rows=3072 width=0) (actual t
("          Index Cond: ((street)::text = 'BROADWAY'::text)",),
('Planning Time: 0.666 ms',),
('Execution Time: 16.249 ms',)]
```

In [83]: %%sql

```
EXPLAIN ANALYSE SELECT * FROM new_york_addresses
WHERE street = '52 STREET';
```

```
* postgresql://postgres:***@localhost:5432/analysis
7 rows affected.
```

```
Out[83]: [('Bitmap Heap Scan on new_york_addresses (cost=5.63..563.21 rows=155 width=46) (actual t
("  Recheck Cond: ((street)::text = '52 STREET'::text)",),
('  Heap Blocks: exact=704',),
('   ->  Bitmap Index Scan on street_idx (cost=0.00..5.59 rows=155 width=0) (actual t
("          Index Cond: ((street)::text = '52 STREET'::text)",),
('Planning Time: 0.163 ms',),
('Execution Time: 5.701 ms',)]
```

In [84]: %%sql

```
EXPLAIN ANALYSE SELECT * FROM new_york_addresses
WHERE street = 'ZWICKY AVENUE';
```

```
* postgresql://postgres:***@localhost:5432/analysis
7 rows affected.
```

```
Out[84]: [('Bitmap Heap Scan on new_york_addresses (cost=5.63..563.21 rows=155 width=46) (actual t
("  Recheck Cond: ((street)::text = 'ZWICKY AVENUE'::text)",),
('  Heap Blocks: exact=6',),
('   ->  Bitmap Index Scan on street_idx (cost=0.00..5.59 rows=155 width=0) (actual t
("          Index Cond: ((street)::text = 'ZWICKY AVENUE'::text)",),
('Planning Time: 0.160 ms',),
('Execution Time: 0.700 ms',)]
```

- we have a parallel seq scan which means that the table will be scanned fully

Adding the Index

In [81]: %%sql

```
CREATE INDEX street_idx ON new_york_adresses (street);

* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out[81]: []

- similar to creating constraints we use **CREATE INDEX** keywords followed by a name we chose for the index (street_idx)
- then **ON** is added following by the target table and column
- this statement will scan the values in the street column and build the index from them
- when we now run the 3 statement we did with the EXPLAIN keyword we can see a better performance
- instead of a seq scan we now have a index scan on street_idx
- instead of visiting each row

– Consider the following two tables from a database you’re making to keep – track of your vinyl LP collection. Start by reviewing these CREATE TABLE – statements.

– The albums table includes information specific to the overall collection – of songs on the disc. The songs table catalogs each track on the album. – Each song has a title and its own artist column, because each song might. – feature its own collection of artists.

```
CREATE TABLE albums ( album_id bigserial, album_catalog_code varchar(100), album_title
text, album_artist text, album_time interval, album_release_date date, album_genre varchar(40),
album_description text );
```

```
CREATE TABLE songs ( song_id bigserial, song_title text, song_artist text, album_id bigint );
```

– Use the tables to answer these questions:

– 1. Modify these CREATE TABLE statements to include primary and foreign keys – plus additional constraints on both tables. Explain why you made your – choices.

In [110]: %%sql

```
CREATE TABLE albums (
    album_id bigserial,
    album_catalog_code varchar(100) NOT NULL,
    album_title text NOT NULL,
    album_artist text NOT NULL,
    album_release_date date,
    album_genre varchar(40),
    album_description text,
    CONSTRAINT album_id_key PRIMARY KEY (album_id),
    CONSTRAINT release_date_check CHECK (album_release_date > '1/1/1925')
);

* postgresql://postgres:***@localhost:5432/analysis
Done.
```

Out[110]: []

In [111]: %%sql

```
CREATE TABLE songs (  
    song_id bigserial,  
    song_title text NOT NULL,  
    song_artist text NOT NULL,  
    album_id bigint REFERENCES albums (album_id),  
    CONSTRAINT song_id_key PRIMARY KEY (song_id)  
);
```

* postgresql://postgres:***@localhost:5432/analysis
Done.

Out[111]: []

- Both tables get a primary key using surrogate key id values that are auto-generated via serial data types.
- The songs table references albums via a foreign key constraint. Note that the reference is a bigint (not a serial anymore)
- In both tables, the title and artist columns cannot be empty, which is specified via a NOT NULL constraint. We assume that every album and song should at minimum have that information.
- In albums, the album_release_date column has a CHECK constraint because it would be likely impossible for us to own an LP made before 1925.

In []: